

How Much of Fake Review Detection on YelpCHI is Business Memorization?

Rithwik Reddy Donthi Reddy
School of Computer Science
University of Oklahoma
Norman, OK, USA
rithwik@ou.edu

Rohan Mukka
School of Computer Science
University of Oklahoma
Norman, OK, USA
rohanmukka@ou.edu

Abstract—YelpCHI is the standard benchmark for fake review detection, and reported accuracies on it are high. We show that much of that performance does not come from detecting deception. Fake reviews in YelpCHI are concentrated by business: 657 of its 1,224 businesses contain no fake reviews at all, and the per-business fake rate is by itself a 0.951-AUC predictor of the label. Any split that lets a business appear in both training and test folds therefore partly measures memorization of which businesses were spam targets. We build ReviewGuard, a detector fusing a lexical text representation with reviewer- and business-level behavioral features, and evaluate it under two protocols differing only in how folds are partitioned. Under a reviewer-disjoint split — the convention in this literature — a behavioral random forest reaches 0.934 AUC-ROC; under a business-disjoint split, with identical features, models, and tuning, it falls to 0.646. Each model’s drop tracks its access to business identity, from 0.032 for a text-only model to 0.289 for the random forest, which is enough to invert the ranking of behavioral and textual features. We additionally document two integrity failures in a widely circulated distribution of the dataset: label-contaminated timestamps yielding 0.965 AUC from the timestamp alone, and a behavioral feature file uncorrelated with both the label and its own column definitions. We argue that business-disjoint evaluation should be reported alongside any YelpCHI result.

Index Terms—fake review detection, evaluation methodology, data leakage, YelpCHI, reviewer behavior, focal loss, multi-modal fusion

I. INTRODUCTION

Online review platforms are consequential enough that manipulating them is profitable. Luca and Zervas estimate that review fraud is economically significant at the scale of entire markets [1], and the availability of fluent text generation has lowered the cost of producing plausible fake content [2]. A large research literature has responded with detectors that combine review text with reviewer behavior, and reported numbers on the standard benchmark — YelpCHI [3] — are strong.

This paper began as an attempt to build such a detector. In the course of doing so we found that the benchmark rewards a shortcut that has little to do with deception, and that the headline numbers a straightforward pipeline produces are largely an artifact of how the data is partitioned.

The mechanism is simple. YelpCHI was assembled by collecting reviews for a set of Chicago businesses. Spam cam-

paigns target particular businesses, so the label is concentrated by business rather than spread uniformly: 657 of the 1,224 businesses in the corpus contain zero fake reviews, the median business has a fake rate of exactly zero, and mapping each review to the fake rate of its business yields an AUC of 0.951 on its own. A model given any feature that identifies the business — its review count, its rating distribution, or a rating deviation computed against its mean — can score well by learning which businesses are spam targets, and will not transfer to a business it has not seen.

We therefore evaluate under two protocols that are identical except for the grouping variable used to form cross-validation folds. The *reviewer-disjoint* protocol guarantees that no reviewer spans the split, which is the precaution the literature normally takes. The *business-disjoint* protocol additionally guarantees that no business spans the split, so a model must generalize to businesses it has never observed. The gap between them measures the shortcut.

Our contributions are:

- 1) A quantification of business-identity leakage in YelpCHI (Section III), and a demonstration that it accounts for the majority of the apparent performance of behavioral fake review detectors on this benchmark (Section VI).
- 2) ReviewGuard, a two-branch detector fusing lexical text features with reviewer and business behavioral features, trained with focal loss under severe class imbalance, evaluated under both protocols with per-model threshold tuning so that baselines are not handicapped (Section IV).
- 3) A data integrity audit of a widely circulated CSV distribution of YelpCHI, identifying label-contaminated timestamps and an uninformative behavioral feature file, both excluded from our experiments (Section III).
- 4) Evidence that the fusion benefit, unlike the behavioral advantage, survives the stricter protocol — alongside the observation that all absolute performance under it is far below what this benchmark is usually reported to support (Section VI).

II. RELATED WORK

A. Detection Methods

Early work treated deceptive review detection as text classification. Ott et al. [4] showed that n -gram and psycholinguistic features separate solicited deceptive reviews from genuine ones with near-human accuracy in a controlled corpus. Jindal and Liu [5] introduced opinion spam as a distinct problem and applied duplicate detection and engineered features at scale. Mukherjee et al. [6] made the central observation that motivates two-branch designs: reviewer behavior — posting bursts, rating extremity, activity patterns — carries signal that text alone does not, and combining the two beats either.

Representation learning followed. Recurrent and convolutional models replaced hand-built lexical features, and pre-trained transformers [7], [8] became the default text encoder, consistently improving on bag-of-words baselines for review classification. Most recently the review ecosystem has been modeled as a heterogeneous graph over reviewers, businesses, and reviews. Rayana and Akoglu [3] propagate spam evidence through reviewer-product networks, and graph neural networks [9], [10] learn representations over the same structure. The YelpCHI distribution used by much of this line of work descends from Rayana and Akoglu’s release.

B. Evaluation and Leakage

Our concern is orthogonal to the choice of model. Kaufman et al. [11] characterize leakage as the introduction of information about the target that would be unavailable at prediction time, and document how easily it survives into published results. Benchmark-specific shortcut learning is now well documented across machine learning [12]: models reach high accuracy through features that correlate with the label in the dataset’s construction rather than in the phenomenon.

Graph-based fake review detectors are structurally exposed to the effect we describe, because the reviewer-business graph makes business identity directly available. This does not invalidate that work — in a deployment where a platform monitors a fixed, known set of businesses, exploiting business-level priors is legitimate. The problem is that the standard random or reviewer-grouped split does not distinguish between a detector that generalizes to new businesses and one that has learned a lookup table of past spam targets, and results are usually not reported in a way that separates the two.

III. DATA INTEGRITY AUDIT

We work from a CSV distribution of YelpCHI containing 45,954 reviews with review text, reviewer and business identifiers, star ratings, timestamps, and labels; 6,677 reviews (14.53%) are labeled fake. The review counts, label counts, and class balance match the canonical .mat release. Before using any column we audited it. Two sources failed, and are excluded from every experiment in this paper. Table I summarizes the diagnostics.

TABLE I

DATA INTEGRITY AUDIT. THE FIRST TWO BLOCKS DESCRIBE INPUTS EXCLUDED FROM THE STUDY; THE THIRD DESCRIBES A PROPERTY OF YELPCHI ITSELF THAT DICTATES THE EVALUATION PROTOCOL.

Source	Diagnostic	Value
date column	AUC from timestamp alone	0.965
	Monthly fake-rate range	0.026–0.498
	Distinct seconds values	1
behavioral_	Max $ r $ vs. own definition	0.005
features.csv	Max $ r $ vs. label	0.008
Business identity	Businesses with no fakes	657/1224
	AUC from business fake rate	0.951

A. Contaminated Timestamps

The released `date` column does not behave like a record of when reviews were posted. A random forest given nothing but the raw Unix timestamp separates fake from genuine reviews with an AUC of 0.965. The monthly fake rate swings from 49.8% in January 2010 to 2.6% in March 2010, and the entire corpus spans 111 days, against the multi-year range of the original Yelp collection. Meanwhile the within-day structure is absent: posting hours are uniform across all 24 hours to within a factor of 1.11, and the seconds field takes exactly one distinct value.

The combination is diagnostic. Real posting times are strongly diurnal, so a uniform hour distribution and a constant seconds field indicate that the time component was generated rather than observed; and a label association this strong within a four-month window indicates that whatever generated the dates was conditioned on the class. Every temporal behavioral feature in the standard repertoire — posting burstiness, inter-review gaps, account age at review — would be computed from this column and would inherit the contamination. We therefore exclude timestamps entirely. This costs us the burstiness features that Mukherjee et al. [6] found informative, and we note it as a limitation rather than a design choice.

B. An Uninformative Feature File

The same distribution ships a six-column feature file, `behavioral_features.csv`, holding average star rating, review count, burst ratio, rating deviation, category diversity, and account age. Reconstructing each column from the review table it purports to summarize — for instance, computing each reviewer’s actual mean rating — yields correlations of at most 0.005 with the shipped values. The columns also correlate with the label at most 0.008 in absolute value, and are near-degenerate: the burst ratio is zero for 97.5% of rows and category diversity is 1.0 for 99.4%. The file carries no recoverable information and we do not use it. We flag it because it is distributed alongside otherwise usable data and its filename invites use without inspection.

C. Business-Concentrated Labels

The third finding is a property of YelpCHI itself rather than a defect in a redistribution, and it is the one that shapes our evaluation. Fake reviews are concentrated by business: 657 of

1,224 businesses contain no fake reviews, 58 consist entirely of fake reviews, and the median business fake rate is zero. Mapping each review to its business’s fake rate gives an AUC of 0.951. Fig. 1 shows the distribution.

This is not an error — it reflects how spam campaigns actually work, and how the corpus was collected. But it means that business identity is close to a sufficient statistic for the label, and that any feature which encodes business identity is a shortcut.

IV. REVIEWGUARD

A. Problem Formulation

Given a review $r = (t, b)$ with text t and behavioral context b , we learn $f : (t, b) \rightarrow [0, 1]$ estimating $P(\text{fake} \mid t, b)$. Positives are reviews labeled as spam in the YelpCHI ground truth.

B. Text Branch

The text branch encodes the review’s actual text. We compute two TF-IDF representations — word unigrams and bigrams, and character n -grams of length three to five with word-boundary padding — each capped at 50,000 features with a minimum document frequency of five and sublinear term frequency scaling. Character n -grams are included because they are robust to the spelling irregularity and punctuation patterns that distinguish review registers. The two matrices are concatenated and reduced by truncated SVD to a 256-dimensional dense vector, which is the representation handed to the fusion model.

To this we append ten stylometric features in the tradition of Ott et al. [4]: character and word counts, mean word length, capitalization, exclamation, question mark, digit and punctuation ratios, first-person pronoun ratio, and type-token ratio. The text branch therefore contributes 266 dimensions.

C. On the Absence of a Transformer Branch

The design we set out to evaluate used a fine-tuned RoBERTa-base encoder [8] in place of the lexical representation, taking the 768-dimensional [CLS] embedding as the text vector. We report results with the lexical encoder instead, for an infrastructural reason we state plainly rather than obscure: the compute environment available for this study had no GPU and no network route to a pre-trained model repository, so the transformer branch could not be run. The encoder is implemented and released with our code so that the comparison can be completed where those resources exist.

This is a real limitation on the text side of our results, and it means our text branch is a lower bound on what a text branch can do. It does not affect the paper’s central claim, which concerns the evaluation protocol and applies to any text encoder: the business-disjoint gap we measure is a property of the benchmark’s label structure, not of the representation used.

D. Behavior Branch

The behavior branch encodes 19 features in three groups. Review-level: star rating, rating extremity $|\text{rating} - 3|$, and review word count. Reviewer-level: review count, a singleton indicator, mean and standard deviation of ratings given, ratio of extreme (one- or five-star) ratings, ratio of positive ratings, number of distinct businesses reviewed, maximum reviews written for any single business, and mean review length. Business-level: review count, mean and standard deviation of ratings received, and the signed and absolute deviation of this reviewer’s rating from its business’s mean.

Two indicator features, `rev_seen_in_training` and `prod_seen_in_training`, record whether the reviewer and business were observed during training, making the cold-start case explicit to the model rather than hiding it in an imputed value.

Crucially, all aggregate tables are fitted on the training rows of a fold and then applied to held-out rows. A reviewer or business absent from training receives training-set priors. Computing these aggregates over the full corpus — the natural and, in our experience, the default implementation — leaks both the existence and the composition of the evaluation fold into the features.

E. Fusion and Training

The 266-dimensional text vector and the 19-dimensional behavior vector are standardized and concatenated, then classified by a two-layer MLP: 256 hidden units, ReLU, dropout 0.3; 64 hidden units, ReLU, dropout 0.3; and a linear output. Single-branch ablations use the same architecture on one branch alone, with a smaller (64, 32) hidden configuration for the behavior-only model to match its input dimensionality.

Training uses focal loss [13],

$$L = -\alpha_t(1 - p_t)^\gamma \log(p_t), \quad (1)$$

with $\gamma = 2$ and α set to the training-set genuine-class frequency, addressing the 14.53% positive rate. We optimize with Adam [14] at a learning rate of 10^{-3} , batch size 256, for at most 40 epochs, retaining the weights with the best validation AUC and stopping after six epochs without improvement.

V. EXPERIMENTAL SETUP

A. Two Protocols

Both protocols use five-fold cross-validation, stratified by label, and differ only in the variable on which folds are made disjoint.

Business-disjoint (primary). Folds are grouped by business identifier, so every business in a test fold is unseen during training. This is the setting in which a detector must generalize to businesses it has no history for.

Reviewer-disjoint (comparison). Folds are grouped by reviewer identifier, so no reviewer spans the split but businesses may. This corresponds to the precaution commonly taken in this literature.

Within each fold we reserve a grouped, stratified validation slice (20% of the training portion, disjoint on the same




Fig. 1. Fake reviews are concentrated by business. Left: distribution of per-business fake rates; the mode at zero is 657 businesses with no fake reviews at all. Right: businesses sorted by fake rate, showing that the label is near-constant within a business.

variable) used for early stopping and threshold selection. Every fitted object — TF-IDF vocabularies, the SVD basis, feature scalers, and the reviewer and business aggregate tables — is fit on the fold’s training portion alone.

B. Threshold Tuning

Under 14.53% positive prevalence, a classifier left at a 0.5 threshold predicts the majority class almost everywhere, and reporting F1 or recall at that threshold says more about calibration than about discrimination. We therefore tune a decision threshold for *every* model, including all baselines, on the validation slice to maximize Macro-F1, and apply it unchanged to the test fold. This matters for fair comparison: in a preliminary version of these experiments, untuned classical baselines recorded fake-class recalls as low as 0.003, which reflects an unadjusted threshold rather than a failure to rank.

C. Baselines

We compare against four baselines and two ablations: TF-IDF with a linear SVM [15] and with logistic regression, both on the sparse lexical matrix; behavioral features with logistic regression and with a random forest [17]; and the text-only and behavior-only MLPs, which serve as the fusion ablations.

D. Metrics and Significance

We report AUC-ROC, average precision, Macro-F1, and fake-class recall, as mean and standard deviation over the five folds. Average precision is included because it is the more informative summary under class imbalance. Pairwise comparisons between ReviewGuard and each baseline use the Wilcoxon signed-rank test over the five paired fold results. We note the limitation this imposes explicitly: with $n = 5$ paired observations, the smallest attainable two-sided p -value is 0.0625, so no comparison can reach $\alpha = 0.05$. We report the p -values for completeness and draw conclusions from effect

sizes and fold-level consistency rather than from significance thresholds.

VI. RESULTS

A. The Protocol Gap

Table II is the paper’s central result. It reports AUC-ROC for every model under both protocols; the models, features, hyperparameters, and tuning procedure are identical, and only the fold partitioning differs.

TABLE II
EFFECT OF THE EVALUATION PROTOCOL ON AUC-ROC. THE REVIEWER-DISJOINT SPLIT ALLOWS A BUSINESS TO APPEAR IN BOTH TRAINING AND TEST FOLDS; THE BUSINESS-DISJOINT SPLIT DOES NOT. THE GAP IS THE PORTION OF APPARENT PERFORMANCE ATTRIBUTABLE TO MEMORISING WHICH BUSINESSES WERE SPAM TARGETS.

Model	Rev.-disj.	Bus.-disj.	Gap
TF-IDF + LinearSVM	0.740	0.666	+0.074
TF-IDF + LogReg	0.762	0.700	+0.062
Behavior + LogReg	0.782	0.599	+0.183
Behavior + RandomForest	0.934	0.646	+0.289
Text-only MLP	0.743	0.712	+0.032
Behavior-only MLP	0.889	0.706	+0.183
ReviewGuard (Fusion)	0.856	0.730	+0.126

The pattern is consistent, interpretable, and ordered exactly by how much business identity each model can reach. The text-only MLP loses 0.032 AUC, because review text identifies its business only weakly. The two TF-IDF baselines lose 0.062 and 0.074. ReviewGuard, which sees business aggregates diluted among 266 text dimensions, loses 0.126. The behavior-only models lose 0.183, and the behavioral random forest — the model best equipped to carve up a small set of business-level aggregates — loses 0.289, falling from the best model in the comparison protocol (0.934) to the worst in the primary one (0.646).




Fig. 2. The same models under both protocols. Bars are means over five folds with standard deviation whiskers. Behavioral models lose most of their apparent performance once businesses no longer span the split; text-based models, which have less access to business identity, lose less.

That reversal is the finding. Under the reviewer-disjoint protocol, the natural reading of Table II is that behavioral features dominate textual ones for fake review detection on YelpCHI, and that a random forest over 19 behavioral features is the model to beat. Under the business-disjoint protocol that conclusion inverts. Most of what the behavior branch contributes on this benchmark is knowing which businesses attract spam, not knowing what a spammer does.

B. Business-Disjoint Performance

Table III reports full metrics under the primary protocol.

ReviewGuard leads on every metric, at 0.730 AUC-ROC and 0.626 Macro-F1. These are modest numbers, and we present them as the honest operating point for this task rather than as a competitive result. Detecting a fake review at a business one has never seen, from text and reviewer statistics alone and without usable timestamps, is genuinely hard. The ordering of the two single branches is also worth noting: text (0.712) slightly ahead of behavior (0.706), the reverse of their ordering under the leakier protocol.

C. Does Fusion Help?

Table IV tests ReviewGuard against each baseline.

Fusion improves on every baseline, and the improvement over the classical models is large: +0.084 AUC over the behavioral random forest and +0.131 over behavioral logistic regression. Against its own ablations the margin is smaller but consistent. ReviewGuard beats the text-only MLP on five of five folds (+0.018 mean AUC) and the behavior-only MLP on four of five (+0.024).

We are deliberate about what this does and does not establish. With five paired folds the Wilcoxon signed-rank test bottoms out at $p = 0.0625$, so the five-of-five comparisons sit at the floor of what the test can report and nothing here clears $\alpha = 0.05$. The effect is also small relative to fold-to-fold spread, which Fig. 4 shows directly. We read the evidence as a real but modest fusion benefit that, unlike the behavioral

advantage, survives the stricter protocol — not as a decisive win. Whether a transformer text branch would widen the margin is exactly the experiment we were unable to run.

D. Discrimination and Errors

Fig. 5 shows ROC curves and Fig. 6 the confusion matrix for ReviewGuard. Both pool the out-of-fold predictions from all five business-disjoint folds, so every review in the corpus is scored exactly once by a model that never saw its business during training.

The error profile is what the threshold objective implies: tuning for Macro-F1 under 14.5% prevalence buys fake-class recall at a substantial cost in false positives. For deployment this trade-off should be set from the relative cost of suppressing a genuine review versus missing a fake one, not from Macro-F1.

E. What the Behavior Branch Actually Uses

The protocol gap in Table II shows that behavioral models rely on business identity; SHAP [16] attribution shows the mechanism directly. Applying a kernel explainer to the behavior-only MLP over a sample of held-out reviews, one feature dominates: `prod_review_count`, the number of reviews the business has received, carries a mean absolute SHAP value of 0.462, five times that of the next feature (`review_word_count`, 0.093). The four features describing the business itself (`prod_*`) account for 55.5% of total attribution while making up 4 of the 19 inputs.

This is worth stating plainly: the most influential input to the behavior branch is a property of the business, not of the review or its author. Under the reviewer-disjoint protocol that feature is highly predictive, because a business seen in training carries its fake rate with it. Under the business-disjoint protocol it is a bare popularity count for a business the model has never seen, and its predictive value largely evaporates. The feature the model leans on hardest is precisely the one that does not transfer.

TABLE III

FIVE-FOLD CROSS-VALIDATED PERFORMANCE ON YELPCHI UNDER THE BUSINESS-DISJOINT PROTOCOL (MEAN $\pm$ S.D. ACROSS FOLDS). DECISION THRESHOLDS FOR EVERY MODEL ARE TUNED ON A HELD-OUT VALIDATION SLICE TO MAXIMISE MACRO-F1.

Model	AUC-ROC	AP	Macro-F1	Rec(Fake)
TF-IDF + LinearSVM	0.666 $\pm$ 0.024	0.265 $\pm$ 0.029	0.587 $\pm$ 0.019	0.299 $\pm$ 0.034
TF-IDF + LogReg	0.700 $\pm$ 0.029	0.305 $\pm$ 0.038	0.608 $\pm$ 0.022	0.361 $\pm$ 0.071
Behavior + LogReg	0.599 $\pm$ 0.020	0.201 $\pm$ 0.010	0.552 $\pm$ 0.007	0.296 $\pm$ 0.120
Behavior + RandomForest	0.646 $\pm$ 0.035	0.221 $\pm$ 0.029	0.567 $\pm$ 0.028	0.286 $\pm$ 0.101
Text-only MLP	0.712 $\pm$ 0.025	0.325 $\pm$ 0.035	0.620 $\pm$ 0.018	0.383 $\pm$ 0.051
Behavior-only MLP	0.706 $\pm$ 0.014	0.295 $\pm$ 0.013	0.612 $\pm$ 0.005	0.375 $\pm$ 0.035
ReviewGuard (Fusion)	0.730 $\pm$ 0.019	0.337 $\pm$ 0.031	0.626 $\pm$ 0.015	0.391 $\pm$ 0.048




Fig. 3. Five-fold performance under the business-disjoint protocol. Error bars are standard deviations across folds.

TABLE IV

WILCOXON SIGNED-RANK TESTS OF REVIEWGUARD AGAINST EACH BASELINE OVER THE 5 BUSINESS-DISJOINT FOLDS. Δ IS REVIEWGUARD'S MEAN PER-FOLD ADVANTAGE. AT $n = 5$ THE SMALLEST ATTAINABLE TWO-SIDED p -VALUE IS 0.0625, SO 0.0625 DENOTES A CLEAN SWEEP OF ALL 5 FOLDS.

Baseline	AUC-ROC		Macro-F1	
	Δ	p	Δ	p
TF-IDF + LinearSVM	+0.064	0.0625	+0.038	0.0625
TF-IDF + LogReg	+0.030	0.0625	+0.018	0.0625
Behavior + LogReg	+0.131	0.0625	+0.074	0.0625
Behavior + RandomForest	+0.084	0.0625	+0.058	0.0625
Text-only MLP	+0.018	0.0625	+0.006	0.1250
Behavior-only MLP	+0.023	0.1250	+0.014	0.1250

VII. DISCUSSION

A. Implications for Benchmarking

We do not claim that prior YelpCHI results are wrong. We claim that the standard protocol does not separate two capabilities that behave very differently in deployment: recognizing deception, and recognizing businesses with a history of being targeted. Both are useful to a platform, and a platform monitoring a fixed set of businesses may legitimately exploit the second. But a detector evaluated only under a business-overlapping split has not been shown to do the first, and its reported number will not survive contact with a new market.




Fig. 4. Per-fold AUC-ROC under the business-disjoint protocol. Horizontal bars mark per-model means. Fold-to-fold spread is comparable to the differences between the leading models, which is why we decline to claim a fusion win.

Reporting both splits costs one extra run and makes the distinction visible.

The point generalizes past this dataset. Any fraud benchmark assembled by sampling entities and collecting their records will concentrate labels by entity, and any model given entity-level aggregates can exploit it.

B. Limitations

Our study has four limitations we consider material.




Fig. 5. ROC curves for ReviewGuard and its two ablations, pooled over the out-of-fold predictions of all five business-disjoint folds.




Fig. 6. ReviewGuard confusion matrix over the pooled out-of-fold predictions of all five business-disjoint folds. A single threshold is fitted to the pooled scores, giving 37.3% fake-class recall against the 39.1% mean of the five per-fold thresholds in Table III.

The text branch is lexical rather than transformer-based, for the infrastructural reason given in Section IV. Our absolute numbers on the text side should be read as a floor.

Timestamps are excluded because they are contaminated in the distribution available to us, which removes the burstiness features that prior work found valuable. A study using a clean temporal record would likely obtain a stronger behavior branch, and should re-examine the protocol gap with those features included — burstiness is plausibly less business-specific than the aggregates we use, and may transfer better.

Five folds admit no significance at $\alpha = 0.05$ under a signed-

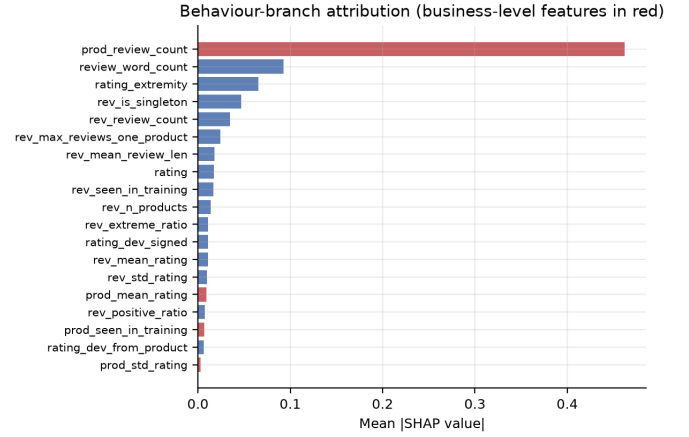


Fig. 7. Mean absolute SHAP value per behavioral feature for the behavior-only MLP. Business-level features are shown in red. The single most influential feature is how many reviews a business has received, which carries no information about the review being classified.

rank test, whose smallest attainable two-sided p -value at $n = 5$ is 0.0625. Our conclusions rest on effect sizes and fold-level consistency, and the fusion comparison in particular is underpowered: a study wanting to settle it should run repeated cross-validation or more folds.

Finally, YelpCHI’s labels originate in Yelp’s proprietary filter. They encode one platform’s operational definition of spam, with unknown false negative rate, and both our results and the leakage we document are properties of that labeling as much as of deception itself.

C. Ethical Considerations

At the operating points reported here, false positives substantially outnumber true positives, and a false positive suppresses a genuine customer’s review. A system of this accuracy is appropriate for ranking a human moderation queue, not for automated removal. We report Macro-F1 and per-class metrics rather than accuracy specifically so that minority-class behavior stays visible.

VIII. CONCLUSION

We set out to build a text-plus-behavior fake review detector for YelpCHI and found that the benchmark’s label structure makes conventional evaluation misleading. Fake reviews cluster by business to the point that business identity alone is a 0.951-AUC predictor, and models evaluated with businesses spanning the split score far higher than the same models evaluated on unseen businesses. The inflation is not uniform: it scales with how directly a model can reach business identity, from 0.032 AUC for a text-only model to 0.289 for a behavioral random forest, which is enough to reverse the ranking of behavioral and textual features. We additionally document label-contaminated timestamps and an uninformative behavioral feature file in a widely circulated distribution of the dataset.

Under business-disjoint evaluation ReviewGuard reaches 0.730 AUC-ROC, ahead of both of its ablations, though by a

margin that five folds cannot establish as significant. We think those numbers, and not the higher ones obtainable under a weaker protocol, are the ones worth reporting.

Three directions follow. First, completing the transformer text branch, which our infrastructure prevented and which is the most likely source of improvement. Second, re-examining published YelpCHI and YelpNYC results under business-disjoint splits, particularly for graph-based methods, which have the most direct access to business identity. Third, obtaining a clean temporal record so that burstiness features can be evaluated for whether they transfer across businesses as poorly as static aggregates do.

ACKNOWLEDGMENT

The authors thank [ADVISOR NAME] for guidance throughout this project.

REPRODUCIBILITY

Code, the data integrity audit, and the metrics behind every table and figure in this paper are available in the project repository. All result tables are generated programmatically from the recorded metrics files rather than transcribed by hand.

REFERENCES

- [1] M. Luca and G. Zervas, “Fake it till you make it: Reputation, competition, and Yelp review fraud,” *Management Science*, vol. 62, no. 12, pp. 3412–3427, 2016.
- [2] E. Crothers, N. Japkowicz, and H. Viktor, “Machine-generated text: A comprehensive survey of threat models and detection methods,” *IEEE Access*, vol. 11, pp. 70977–71002, 2023.
- [3] S. Rayana and L. Akoglu, “Collective opinion spam detection: Bridging review networks and metadata,” in *Proc. ACM KDD*, Sydney, Australia, 2015, pp. 985–994.
- [4] M. Ott, Y. Choi, C. Cardie, and J. T. Hancock, “Finding deceptive opinion spam by any stretch of the imagination,” in *Proc. ACL*, Portland, OR, 2011, pp. 309–319.
- [5] N. Jindal and B. Liu, “Opinion spam and analysis,” in *Proc. ACM WSDM*, Palo Alto, CA, 2008, pp. 219–230.
- [6] A. Mukherjee, V. Venkataraman, B. Liu, and N. S. Glance, “What Yelp fake review filter might be doing?,” in *Proc. ICWSM*, Cambridge, MA, 2013, pp. 409–418.
- [7] J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova, “BERT: Pre-training of deep bidirectional transformers for language understanding,” in *Proc. NAACL-HLT*, Minneapolis, MN, 2019, pp. 4171–4186.
- [8] Y. Liu, M. Ott, N. Goyal, J. Du, M. Joshi, D. Chen, O. Levy, M. Lewis, L. Zettlemoyer, and V. Stoyanov, “RoBERTa: A robustly optimized BERT pretraining approach,” arXiv:1907.11692, 2019.
- [9] T. N. Kipf and M. Welling, “Semi-supervised classification with graph convolutional networks,” in *Proc. ICLR*, Toulon, France, 2017.
- [10] Z. Wang, K. Hou, D. Song, Z. Li, T. Zhang, and G. Hu, “Attributed graph neural network for fake review detection,” in *Proc. IEEE ICDM*, Sorrento, Italy, 2020, pp. 621–630.
- [11] S. Kaufman, S. Rosset, C. Perlich, and O. Stitelman, “Leakage in data mining: Formulation, detection, and avoidance,” *ACM Trans. Knowledge Discovery from Data*, vol. 6, no. 4, pp. 1–21, 2012.
- [12] R. Geirhos, J.-H. Jacobsen, C. Michaelis, R. Zemel, W. Brendel, M. Bethge, and F. A. Wichmann, “Shortcut learning in deep neural networks,” *Nature Machine Intelligence*, vol. 2, no. 11, pp. 665–673, 2020.
- [13] T.-Y. Lin, P. Goyal, R. Girshick, K. He, and P. Dollár, “Focal loss for dense object detection,” in *Proc. IEEE ICCV*, Venice, Italy, 2017, pp. 2980–2988.
- [14] D. P. Kingma and J. Ba, “Adam: A method for stochastic optimization,” in *Proc. ICLR*, San Diego, CA, 2015.
- [15] C. Cortes and V. Vapnik, “Support-vector networks,” *Machine Learning*, vol. 20, no. 3, pp. 273–297, 1995.
- [16] S. M. Lundberg and S.-I. Lee, “A unified approach to interpreting model predictions,” in *Proc. NeurIPS*, Long Beach, CA, 2017, pp. 4765–4774.
- [17] L. Breiman, “Random forests,” *Machine Learning*, vol. 45, no. 1, pp. 5–32, 2001.